

01 k8s basics and architecture

Container Orchestration and Kubernetes

Why Container Orchestration?

Containerization (e.g., using Docker) packages an application and all its dependencies into a single, portable unit. While managing a single container is simple, modern applications are often **distributed systems** composed of **microservices**, requiring dozens, hundreds, or even thousands of containers.

Managing this scale manually is practically impossible and introduces significant challenges.

Container orchestration is the process of automating the deployment, scaling, management, networking, and availability of containerized applications across a cluster of hosts.

Challenges Solved by Orchestration:

- **Scaling and Load Balancing:** How do you run multiple copies (replicas) of a container to handle high traffic, distribute incoming network requests across them, and scale them up/down automatically based on demand?
- **High Availability and Self-Healing:** What happens if a server (node) fails, or a container crashes? An orchestration tool must detect the failure and automatically restart the container or reschedule it on a healthy node to maintain application uptime.
- **Networking and Service Discovery:** Containers are ephemeral and their IP addresses change. How do different microservices find and communicate with each other without hardcoding IP addresses?
- **Deployment Management:** How do you update an application to a new version without downtime (**rolling updates**) and safely roll back if the new version has a critical bug?
- **Resource Optimization:** How do you efficiently pack containers onto the available hosts (nodes) to maximize hardware utilization and minimize costs?
- **Configuration and Secret Management:** How do you inject configuration data (e.g., database URLs, API keys) into containers securely and consistently across different environments?

Container orchestration platforms provide a **declarative approach**—you define the *desired state* of your application (e.g., "I want 3 replicas of my web app running"), and the orchestrator is responsible for moving the *actual state* to match the desired state, continually.

Introduction to Kubernetes and its Architecture

Kubernetes (K8s) is an open-source system for automating deployment, scaling, and management of containerized applications. Originally designed by Google, it is now maintained by the Cloud Native Computing Foundation (CNCF).

The Core Concept: The Cluster

A Kubernetes cluster is a set of machines (physical or virtual) that run containerized applications. The cluster is divided into two main parts:

1. **The Control Plane (Master Node):** The "brain" of the cluster that manages the overall state and makes global decisions (like scheduling).
2. **Worker Nodes:** The machines that run the actual application workloads (containers) in the form of **Pods**.

Control Plane Components

These components run on the master node(s) and manage the cluster's state:

- **API Server (`kube-apiserver`):**
 - The **front-end** for the Control Plane. It exposes the Kubernetes API.
 - It is the only component that communicates with the cluster's persistent store. All internal and external requests (e.g., from `kubectl`) go through the API Server.
- **etcd:**
 - A **consistent and highly available key-value store** that holds the entire cluster state, configuration data, and metadata.
 - It is the single source of truth for the cluster.
- **Scheduler (`kube-scheduler`):**
 - Watches for newly created **Pods** that have no node assigned.
 - Selects an optimal node for the Pod to run on based on various factors like resource requirements (CPU/Memory), resource availability, policy constraints, and affinity/anti-affinity specifications.
- **Controller Manager (`kube-controller-manager`):**
 - Runs controller loops that regulate the state of the cluster.
 - Each controller attempts to move the current state closer to the desired state. Examples include:
 - **Replication Controller:** Ensures the specified number of Pod replicas is running.
 - **Node Controller:** Notices when a node goes down and handles its status.
 - **Endpoint Controller:** Populates the Endpoints object (used for Service discovery).

Worker Node Components


These components run on every worker node and are responsible for running and managing containers:


- **Kubelet:**
 - The **primary agent** that runs on each node.
 - It communicates with the Control Plane and ensures that containers described in **Pods** are running and healthy on its node.
 - It reports the node's health and resource utilization back to the master.
- **Container Runtime:**
 - The software responsible for running the containers (e.g., Docker, containerd, CRI-O).

- The Kubelet interacts with the Container Runtime to start, stop, and manage container lifecycles.
- **Kube-Proxy (kube-proxy):**
 - A network proxy that runs on each node.
 - It maintains network rules on the nodes, allowing network communication to your Pods from outside and inside the cluster (e.g., implementing **Services**).


Key Kubernetes Objects (Abstractions)

- **Pod:** The **smallest deployable unit** in Kubernetes. It is a group of one or more containers that are deployed together on the same host and share network and storage resources. A Pod represents a single instance of an application.
- **Deployment:** A declarative way to manage the creation and scaling of **Pods** and **ReplicaSets**. It handles rolling updates and rollbacks, allowing for zero-downtime application updates.
- **Service:** An abstraction that defines a logical set of **Pods** and a policy by which to access them. Services provide a **stable IP address and DNS name** for a set of Pods, enabling internal service-to-service communication and external exposure of the application.





Kubernetes Architecture



Kubernetes vs Docker Swarm vs Other Platforms

Feature	Kubernetes (K8s)	Docker Swarm	Apache Mesos (with Marathon)
Complexity	High. Steeper learning curve, more components to manage.	Low. Simpler to set up, highly integrated with the Docker ecosystem.	Medium/High. Designed to manage a data center's resources, not just containers.
Ecosystem & Community	Massive. The industry standard, vast tooling, and active community.	Moderate. Good integration with Docker tools, but smaller ecosystem for advanced features.	Moderate/Niche. Powerful, but often used for diverse workloads (containers, Big Data, etc.).
Scalability (Max Nodes)	High. Can scale to thousands of nodes, designed for massive scale.	Medium. Suitable for small to medium clusters.	Extremely High. Designed to scale to tens of thousands of nodes (Data Center OS).
Networking	Advanced. Supports a wide range of networking models (e.g., CNI), and sophisticated Service objects.	Simple. Built-in overlay networking, simpler load balancing.	Customizable. Highly flexible resource management.

Feature	Kubernetes (K8s)	Docker Swarm	Apache Mesos (with Marathon)
Auto-Scaling	Native. Built-in Horizontal Pod Autoscaler (HPA) and Cluster Autoscaler.	Limited/External. Requires manual scaling or third-party tools/scripts.	Framework-dependent. Managed via frameworks like Marathon.
Use Case	Complex, large-scale microservices, fluctuating workloads, high customization needs.	Quick deployment, small to medium-sized clusters, simplicity is key, existing Docker user base.	Environments with diverse workloads (Big Data, HPC, containers) and extreme fault-tolerance needs.

When to Use Kubernetes

While powerful, Kubernetes adds an operational overhead. It is a good fit when the added complexity is justified by the application's requirements.

Use Cases Where Kubernetes Excels

- **Microservices Architecture:** When your application is composed of many independent services that need advanced service discovery, load balancing, and independent scaling.
- **High Availability (HA) and Resilience:** For mission-critical applications where downtime must be minimized. K8s' self-healing capabilities are crucial for ensuring application continuity against node or container failure.
- **Complex CI/CD Pipelines:** When you need a reliable, repeatable, and automated deployment platform that supports rolling updates, canary deployments, and blue/green deployments across different environments (Dev, Test, Prod).
- **Fluctuating Workloads:** For applications with unpredictable or periodic traffic spikes (e.g., e-commerce, seasonal events). The **Horizontal Pod Autoscaler (HPA)** can automatically scale up and down to match demand and optimize cloud costs.
- **Cloud and Vendor Portability:** When you need to ensure your application can run consistently across different public cloud providers (AWS, Azure, GCP) or on-premises data centers without significant changes.
- **Multi-Tenancy:** When running applications for multiple customers (tenants) on the same cluster, leveraging **Namespaces** to provide logical isolation and resource quotas.

When to Avoid or Delay Kubernetes

- **Simple, Small Applications:** For a small application (e.g., a single-container website) that doesn't need high availability or complex scaling, the operational overhead of managing a K8s cluster is generally not worth the benefit. Simple solutions like **Docker Compose** or managed services (e.g., AWS Fargate, Google Cloud Run) are better choices.
- **Small Teams/Limited Expertise:** If your team lacks the necessary skills, adopting Kubernetes can significantly slow down development due to the steep learning curve and the complexity of debugging cluster-level issues.

- **Monolithic Applications with No Scaling Needs:** If you have a traditional application that is not broken into microservices and can handle its load on one or two VMs, a traditional deployment model will be simpler and more cost-effective.

<https://aws.plainenglish.io/kubernetes-architecture-c93cb9c798d8>