

## 04 K8s Service 101

### Kubernetes Services

#### Overview

A Kubernetes Service is a logical abstraction that provides a stable network endpoint for accessing a set of Pods. Services enable reliable service discovery and load distribution within a Kubernetes cluster by creating a consistent interface to dynamic Pod resources.

#### Core Concepts

##### Service Fundamentals

A Service acts as an internal proxy layer that:

- Provides a stable IP address and DNS name for Pod access
- Automatically discovers and tracks Pods matching specified label selectors
- Distributes incoming traffic across healthy Pod instances
- Maintains service continuity as Pods are created, terminated, or rescheduled

Unlike traditional load balancers, Kubernetes Services operate through iptables rules managed by kube-proxy, consuming minimal cluster resources while providing essential traffic routing capabilities.

##### Service Architecture

Services decouple service consumers from the underlying Pod infrastructure. When a client application needs to communicate with another service:

1. The client references the Service by name (e.g., `http://product-service`)
2. Cluster DNS resolves the Service name to its stable ClusterIP
3. kube-proxy maintains iptables rules routing traffic to appropriate Pod IPs
4. The Service distributes requests to healthy Pods matching its selector

This architecture eliminates the need for clients to track individual Pod IP addresses or manage service discovery manually.

## Service Specification

### Basic Service Definition

```
apiVersion: v1
kind: Service
metadata:
  name: nginx
spec:
  selector:
    app: my-app
  ports:
  - port: 80
    targetPort: 80
    protocol: TCP
```

YAML

## Specification Components

### Metadata

- **name**: Service identifier used by clients for DNS-based discovery

### Selector

- Defines label criteria for Pod selection
- Services route traffic only to Pods matching all specified labels
- Multiple labels enable precise Pod targeting in large clusters

### Ports Configuration

- **port**: The port exposed by the Service (client-facing port)
- **targetPort**: The port on the Pod containers receiving traffic
- **protocol**: Network protocol (TCP/UDP/SCTP), defaults to TCP

The **port** and **targetPort** serve distinct purposes: clients connect to the Service's **port**, while the Service forwards traffic to the Pod's **targetPort**. These values may differ to accommodate port standardization or security requirements.

## Service Discovery

### DNS-Based Discovery

Kubernetes runs an internal DNS service that automatically creates DNS records for all Services. Applications can reference Services using their names rather than IP addresses:

```
curl http://product-service
```

SHELL

The DNS service resolves the name to the Service's ClusterIP, enabling location-independent service access.

## Endpoint Management

Services maintain an endpoints list containing IP addresses of all matching Pods. The Kubernetes control plane continuously updates these endpoints as Pods are created or terminated, typically within milliseconds of Pod state changes.

You can inspect Service endpoints using:

```
kubectl describe service <service-name>
```

SHELL

## kube-proxy and Traffic Routing

### Traffic Distribution Mechanism

The kube-proxy component runs on every cluster node and maintains iptables rules that implement Service routing. When traffic arrives at a Service IP:

1. iptables rules intercept the traffic
2. kube-proxy selects a backend Pod (not strictly round-robin)
3. Traffic is forwarded directly to the chosen Pod's IP address

This implementation means Services themselves don't consume CPU or memory—they exist purely as networking configurations.

## Load Distribution Characteristics

Services perform random load distribution rather than traditional round-robin scheduling. The algorithm:

- Distributes traffic across all healthy Pod endpoints
- Does not maintain session affinity by default
- Routes only to Pods passing readiness checks

Services are not designed for sophisticated traffic management (path-based routing, weighted distribution, etc.). For advanced routing requirements, consider using Ingress controllers or dedicated load balancing solutions.

## Service Types

Kubernetes provides three primary Service types, each suited for different networking requirements.

### ClusterIP (Default)

#### **Characteristics:**

- Provides internal-only cluster networking

- Accessible exclusively within the Kubernetes cluster
- Default type when no type is specified
- Ideal for service-to-service communication in microservices architectures

### **Use Cases:**

- Backend services not requiring external access
- Database connections
- Internal APIs
- Cache services (Redis, Memcached)

### **Configuration:**

```
apiVersion: v1
kind: Service
metadata:
  name: redis
spec:
  type: ClusterIP # Optional, as this is default
  selector:
    app: redis
  ports:
  - port: 6379
    targetPort: 6379
```

YAML

## **NodePort**

### **Characteristics:**

- Extends ClusterIP functionality with external accessibility
- Opens a specific port (30000-32767 range) on every cluster node
- Requests to any node's NodePort are forwarded to the Service
- Primarily used for development and testing scenarios

### **Use Cases:**

- Local development testing
- Temporary external access requirements
- Testing service configurations before production deployment

### **Configuration:**

```

apiVersion: v1
kind: Service
metadata:
  name: my-app
spec:
  type: NodePort
  selector:
    app: my-app
  ports:
  - port: 80
    targetPort: 80
    nodePort: 30001 # Optional: specify exact port, otherwise randomly assigned

```

### Important Considerations:

- NodePort services remain accessible as ClusterIP services internally
- All cluster nodes listen on the specified port
- Not recommended for production external traffic (use Ingress instead)
- Requires firewall configurations in cloud environments

## LoadBalancer

### Characteristics:

- Automatically provisions cloud provider load balancers
- Available only in cloud-managed Kubernetes environments (GKE, EKS, AKS)
- Creates external IP addresses for public access
- Builds upon NodePort functionality

### Cloud Provider Integration:

- AWS: Creates Classic Load Balancer or Network Load Balancer
- GCP: Creates TCP/UDP Load Balancer
- Azure: Creates Azure Load Balancer

### Configuration:

```
apiVersion: v1
kind: Service
metadata:
  name: web-service
spec:
  type: LoadBalancer
  selector:
    app: web
  ports:
  - port: 80
    targetPort: 8080
```

### Limitations:

- Not functional in local development environments (kind, minikube)
- Each Service creates a separate cloud load balancer (can incur costs)
- For production HTTP/HTTPS traffic, Ingress resources are typically preferred

## Practical Implementation Patterns

### Multi-Tier Application Architecture

For applications with multiple components (application tier, cache layer, database):

```
# Application Deployment
apiVersion: apps/v1
kind: Deployment
metadata:
  name: app
spec:
  replicas: 3
  selector:
    matchLabels:
      app: application
  template:
    metadata:
      labels:
        app: application
    spec:
      containers:
        - name: app
          image: myapp:v1
          ports:
            - containerPort: 8080
          env:
            - name: REDIS_HOST
              value: "redis"
            - name: REDIS_PORT
              value: "6379"
---
# Redis Deployment
apiVersion: apps/v1
kind: Deployment
metadata:
  name: redis
spec:
  replicas: 1
  selector:
    matchLabels:
      app: redis
  template:
    metadata:
      labels:
        app: redis
    spec:
      containers:
        - name: redis
          image: redis:6
          ports:
```

```

- containerPort: 6379
---

# Redis Service
apiVersion: v1
kind: Service
metadata:
  name: redis
spec:
  selector:
    app: redis
  ports:
  - port: 6379
    targetPort: 6379

```

In this pattern, the application references Redis using the service name `redis`, which resolves to the Redis Service's ClusterIP. This approach eliminates hardcoded IP addresses and enables seamless component scaling.

## Service Testing and Debugging

For testing ClusterIP services that aren't externally accessible, deploy a utility Pod within the cluster:

```

apiVersion: v1
kind: Pod
metadata:
  name: debug-pod
spec:
  containers:
  - name: debug
    image: curlimages/curl
    command: ["sleep", "3600"]

```

YAML

Access the Pod and test services:

```

kubectl exec -it debug-pod -- sh
curl http://my-service

```

SHELL

## Load Distribution Verification

To verify Service load distribution across multiple replicas:

```

for i in {1..1000}; do
  curl -s http://my-app | grep -o "<title>[^<]*" | tail -c+8
done

```

SHELL

This script sends multiple requests and displays which Pod handled each request, demonstrating the Service's load distribution behavior.

## Service vs. Traditional Service Discovery

Kubernetes Services provide built-in service discovery, potentially eliminating the need for dedicated service discovery platforms (Consul, Eureka, etc.) in Kubernetes environments.

### Advantages of Kubernetes Services:

- Zero additional infrastructure or maintenance
- Automatic integration with all cluster components
- Native Kubernetes resource management
- DNS-based discovery without client libraries

### When to Consider External Service Discovery:

- Cross-cluster service discovery requirements
- Advanced service mesh features
- Legacy application compatibility
- Specific organizational standards

## Best Practices

1. **Use ClusterIP by default:** Most services should use ClusterIP unless external access is required
2. **Leverage DNS names:** Reference services by name rather than IP addresses for maintainability
3. **Implement proper selectors:** Use multiple labels to precisely target Pods and avoid unintended routing
4. **Standardize ports:** Align Service ports with container ports unless specific requirements dictate otherwise
5. **Document port mappings:** Clearly document the relationship between Service ports and container ports
6. **Use Ingress for HTTP/HTTPS:** For production external access, use Ingress resources rather than NodePort or LoadBalancer Services
7. **Monitor endpoint health:** Implement readiness probes to ensure Services route only to healthy Pods
8. **Consider session affinity:** For stateful applications, configure sessionAffinity if needed

## Advanced Considerations

### Session Affinity

By default, Services don't maintain session affinity. To enable client IP-based session affinity:

YAML

```

spec:
  sessionAffinity: ClientIP
  sessionAffinityConfig:
    clientIP:
      timeoutSeconds: 3600

```

## Headless Services

For scenarios requiring direct Pod access without load balancing:

YAML

```

spec:
  clusterIP: None
  selector:
    app: my-app

```

Headless Services return Pod IPs directly via DNS rather than a single Service IP.

## External Services

Services can also represent external resources outside the cluster using ExternalName or explicit endpoints, enabling consistent service discovery patterns across internal and external dependencies.

## Integration with Ingress

While NodePort and LoadBalancer Services provide external access, production HTTP/HTTPS applications typically use Ingress resources, which offer:

- Path-based routing
- Host-based routing
- TLS termination
- Single external IP for multiple services
- Advanced traffic management

Services act as backends for Ingress resources, with Ingress controllers managing external traffic routing to appropriate Services.

## Conclusion

Kubernetes Services provide essential abstraction for reliable, scalable service communication within clusters. By understanding Service types, traffic routing mechanisms, and implementation patterns, you can design robust microservices architectures that leverage Kubernetes' built-in networking capabilities effectively. Services eliminate the complexity of manual service discovery and load balancing, enabling teams to focus on application logic rather than infrastructure management.